

CS3103: Operating Systems (Spring 2019)

Programming Assignment 3 P3: A Simulator for FIFO Queue with Progressive Flow

1 Goals

This assignment is designed to help you:

1. get familiar with multi-threaded programming with Pthread.
2. get familiar with critical section with Pthread mutex.
3. get familiar with synchronization with Pthread semaphore. Note that **any semaphore should be given an initial value, or otherwise it is logically wrong!**

You are required to implement your solution in C++ (other languages are not allowed). Your work will be tested on the Linux server cs3103-01.cs.cityu.edu.hk.

A good tutorial on pthread could be found at <https://computing.llnl.gov/tutorials/pthreads/>.

2 Requirements

2.1 Background

This assignment is built above your solution to Assignment 2. To make your life easy, you should re-use your code in Assignment 2.

In this assignment, we add a new type of flow, called progressive flow (**pflow** for short), to the system described in Assignment 2.

The behavior of **pflow** is as follows: **whenever the queue is emptied by the server¹**, it generates a random number of tokens, uniformly distributed in $[1, 5]$, and inserts the tokens to the queue. Otherwise, it is blocked if the queue is not empty.

The behavior of **flow**, **queue**, and **server** are the same as in Assignment 2. Note that the sequence number of each token should be unique and in an increasing order in the system. Same as in Assignment 2, the simulation stops after **MaxToken** number of tokens have been served, where **MaxToken** is an input parameter.

¹For simplicity, we ignore the situation that **pflow** runs before **flow** has injected some tokens in the queue.

2.2 Implementation Requirement

Your program, named by `PFIFO`, should consist of four threads: the main thread, the flow thread, the pflow thread, and the server thread. The main thread is responsible for creating the other three threads, and then waits for them to return. The flow thread simulates the behavior of the flow. The pflow thread simulates the behavior of the pflow. The server thread simulates the behavior of the server.

Same as in Assignment 2, the queue is implemented as a data structure, rather than a thread. You need to implement **your own queue data structure and its related functions**, which means that **you cannot use the queue data structure provided by the C/C++ library**.

Since `flow`, `pflow`, and `server` all need to access the queue data structure, the code segment for accessing the queue is a critical section and must be protected with the pthread mutex lock.

Since the sequence number of each token should be unique and in an increasing order in the system, you need to introduce a global variable `currentSeq` to tell the current sequence number to `flow` and `pflow` whenever they are active. There will be a race condition on accessing `currentSeq` between `flow` and `pflow`, and hence the code for accessing `currentSeq` should be protected with a pthread mutex lock.

Note that there is a synchronization issue between the `pflow` and the `server`: When the `server` empties the queue, it must notify the `pflow` to insert tokens into the queue. However, if the `server` empties the queue **again** before the `pflow` responds to the last notification, the `server` **must not** repeat the notification. You need to use semaphore to fulfill such synchronization.

Same as in Assignment 2, your code **must** print out useful information, and the output **must** be formatted as the following table:

Flow		Queue		Server	
Token added	Latest sequence number	Current	Length	Token fetched	Total Token fetched
aaa(flow)	bbb	ccc			
		zzz		xxx	yyy

where the header of the table is printed by the main thread before creating other threads. The numbers `aaa`, `bbb`, and `ccc` are printed by the `flow` (or `pflow`) when it wakes up. You must label after `aaa` with either `(flow)` or `(pflow)` to show the creator of the tokens. The numbers `xxx`, `yyy`, and `zzz` are printed by the `server` when it wakes up.

In addition, the number `aaa` is **not** an accumulative number but the number of tokens at the current round of flow/pflow wake-up; the number `bbb` is the highest sequence number at the current round of flow/pflow wake-up; the number `ccc` is the queue length after tokens generated at the current round of flow/pflow wake-up have been added into the queue. The number `xxx` is **not** an accumulative number but the number of tokens that will be fetched by the server at the current round of its wake-up; the number `yyy` is the total number of tokens that have been fetched by the server; the number `zzz` is the current queue length after tokens have been fetched by the server at the current round of server wake-up.

Note that at the end of the program, the total number of tokens that have been served is equal to the total number of tokens that have been fetched by the server and the total number of tokens that have been dropped by the queue. As such, **at the end of the program**, you must print out the total number of tokens that have been fetched by the server thread, the total number of tokens that have been generated by the flow thread, the total number of tokens that have been generated

67 by the pflow thread, and the total number of tokens that have been dropped by the queue. The
68 output should be:

69 The total number of tokens that have been fetched by the server is eee.
70 The total number of tokens that have been generated by the flow is fff.
71 The total number of tokens that have been generated by the pflow is ggg.
72 The total number of tokens that have been dropped by the queue is hhh.

73 where eee, fff, ggg, and hhh indicate the corresponding numbers.

74 **Remark 1:** After the pflow is woken up by the server, there might be a competition between
75 the flow and the pflow for accessing the queue. If the flow wins the competition and inputs the
76 tokens before the pflow, the pflow will inject tokens into a non-empty queue. Nevertheless, this
77 special case is acceptable.

78 **Remark 2:** Since the number of last-batch tokens fetched by the server before it stops is
79 random, the sum of eee and hhh might be equal to or slightly larger than MaxToken. Both cases
80 are acceptable. This relaxation is to ease your coding.

81 **Remark 3:** Refer to the flowchart for a baseline implementation.

82 2.3 Prompt for user input

83 Your PFIFO needs to accept two integer inputs, MaxToken as the first and flowInterval as the
84 second, e.g.,

85 \$./PFIFO 500 2

86 2.4 Code Quality

87 We cannot specify completely the coding style that we would like to see but it includes the following:

- 88 1. Proper decomposition of a program into subroutines (and multiple source code files when
89 necessary)—A 500 line program as a single routine won't suffice.
- 90 2. Comment—judiciously, but not profusely. Comments also serve to help a marker, in addition
91 to yourself. To further elaborate:
 - 92 (a) Your favorite quote from Star Wars or Douglas Adams' Hitch-hiker's Guide to the Galaxy
93 does not count as comments. In fact, they simply count as anti-comments, and will result
94 in a loss of marks.
 - 95 (b) Comment your code in English. It is the official language of this university.
- 96 3. Proper variable names—leia is not a good variable name, it never was and never will be.
- 97 4. Small number of global variables, if any. Most programs need a very small number of global
98 variables, if any. (If you have a global variable named temp, think again.)
- 99 5. The return values from all system calls and function calls listed in the assignment
100 specification should be checked and all values should be dealt with appropriately.

3 Marking

Your program will be tested on our CSLab Linux servers (cs3103-01, cs3103-02, cs3103-03).

You should tell TA how to compile and run your code in your [readme](#) text file.

TAs are not supposed to fix the bugs, either in your source code or in your [Makefile](#).

If an executable file cannot be generated and running successfully on our Linux servers, it will be considered as unsuccessful.

Any program does not print the output as required in Section 2.2 will be considered as unsuccessful.

If the program can be run successfully and output is in the correct form, your program will be graded with the following marking scheme.

Table 1: Marking scheme.

Components	Weight
Initialization of semaphore	5%
Correct thread create/join/exit	15%
Correctness of programming logic	10%
Correct use of semaphore	20%
Synchronization with correct mutex	15%
Correctness of printed summary	10%
Queue structure & related functions	10%
Error handling	10%
Programming style and in-program comments	5%

4 Submission

1. This assignment is to be done individually or by a group of two students. You are encouraged to discuss the high-level design of your solution with your classmates but you must implement the program on your own. Academic dishonesty such as copying another student's work or allowing another student to copy your work, is regarded as a serious academic offence.
2. Each submission consists of three files: a source program file (.cpp file), a readme (.txt) file telling us how to compile your code, a text file containing the possible outputs produced by your program (.txt file), and a [Makefile](#) if applicable.
3. Use your student ID(s) to name your submitted files, such as 5xxxxxxx.cpp, 5xxxxxxx-readme.txt, 5xxxxxxx.txt for individual submission, or 5xxxxxxx_5yyyyyyy.cpp, 5xxxxxxx_5yyyyyyy-readme.txt, 5xxxxxxx_5yyyyyyy.txt for group submission. Only ONE submission is required for each group.
4. Submit the files to Canvas.
5. The deadline is 10:00 am, 11-APR. No late submission will be accepted.

Question?

Contact Mr LIU, Dawei at daweiliu2-c@my.cityu.edu.hk or your course lecturer.

The End